Qwen

The Mythos Moment: Assessing the Real Impact of Claude's Frontier AI on Cybersecurity Strategy and Enterprise Risk

Defining the Disruption: An Overview of Claude's Mythos Model

Anthropic's 'Claude Mythos' represents a significant development in the field of artificial intelligence, specifically positioned as a frontier model with advanced capabilities relevant to complex, high-stakes domains [8]. Its emergence has generated considerable discourse within the cybersecurity industry, framed by some as a paradigm-shifting technology and by others as a calculated marketing initiative. Understanding the precise nature of Mythos, its intended purpose, and the context of its introduction is fundamental to assessing its true impact. The model was officially launched as a gated research preview on April 7, 2026, with a knowledge cutoff date of December 2025 [8]. This phased release strategy, granting access prioritized for defensive cybersecurity use cases, underscores its experimental nature and Anthropic's focus on responsible deployment through initiatives like Project Glasswing [1] [8].

The most prominent technical differentiator of Mythos Preview is its exceptionally large context window, which supports up to 1 million tokens [8]. This capability allows the model to process and reason over vast amounts of information simultaneously, such as entire source code repositories, complex network architectures, or extensive historical data sets, without losing contextual integrity. This contrasts sharply with many general-purpose LLMs that operate with much smaller context windows, often limiting their ability to perform holistic analysis of complex systems. The model also features a maximum output length of 128,000 tokens and is explicitly built for reasoning tasks [8]. It is available via the `anthropic.claude-mythos-preview` model ID on the Amazon Bedrock platform, indicating a cloud-based, scalable delivery mechanism designed for enterprise integration [8]. Unlike the more widely available Claude Opus, Sonnet, and Haiku models, Mythos is not listed on Anthropic's primary public models overview page,

reinforcing its status as a specialized, research-focused tool rather than a flagship consumer product [3].

The core design philosophy behind Mythos appears to be centered on creating powerful, autonomous agents capable of executing ambitious, long-running projects [8]. According to CrowdStrike's assessment, the model expands the reasoning, planning, and execution capabilities of AI agents, which directly impacts the endpoint where data access, decisions, and risk are born [1]. This focus on agentic behavior—where an AI can autonomously plan and execute multi-step tasks—is a critical feature distinguishing it from models primarily used for generating text or images in response to discrete prompts. The model's development under Anthropic's Responsible Scaling Policy further indicates an attempt to manage the inherent risks associated with such powerful technology [1]. However, the official naming convention adds a layer of complexity to its identity. Within the leaked documents and subsequent industry discussion, the model was referred to as 'Mythos', but another alias, 'Claude Capybara', was also used [2]. This dual nomenclature suggests internal project codenames may have circulated before the official launch, potentially contributing to the confusion surrounding its exact specifications and capabilities at the time of its leak.

The initial market reaction to the concept of Mythos was swift and significant. A leak revealing Anthropic was testing an advanced AI model with superior coding and threat detection capabilities triggered a sharp sell-off in cybersecurity stocks [4]. This demonstrates that investors perceive the technology as a major disruptive force, capable of reshaping the competitive dynamics and value proposition of the security industry. However, alongside this market-driven anxiety, there exists a strong current of skepticism within the broader cybersecurity community. Some professionals view the hype surrounding Mythos as a sophisticated marketing campaign or a public relations effort by Anthropic, possibly aimed at securing additional funding from investors during a period of intense competition [5]. This skepticism is rooted in several factors. First, the official Anthropic website does not feature a model named 'Mythos', instead listing the latest generation as Claude Opus 4.6, Claude Sonnet 4.6, and Claude Haiku 4.5 [3]. This discrepancy between leaked information and official communications fuels doubts about whether Mythos is a genuinely novel breakthrough or simply a rebranded version of existing models. Second, the rapid pace of AI development means that claims of superiority can quickly become outdated, leading to a degree of wariness among practitioners accustomed to seeing incremental advancements presented as revolutionary [5]. Finally, the rise of powerful open-weight models that can run locally on consumer-grade hardware challenges the premise that only proprietary, expensive systems are required for advanced security tasks, suggesting that the tools for sophisticated analysis

are becoming increasingly democratized [5]. Therefore, while the leaked details point to a technologically impressive model with unique capabilities, its ultimate impact will depend not only on its performance but also on how it is perceived and adopted by the market, balancing genuine innovation against potential hype.

Feature	Claude Mythos Preview	Comparable Models (Opus/Sonnet/Haiku)
Official Name/ Listing	Not listed; referred to as 'Mythos' or 'Claude Capybara' in leaks [2] [3]	Listed as Claude Opus 4.6, Claude Sonnet 4.6, and Claude Haiku 4.5 [3]
Primary Use Case	Ambitious projects in cybersecurity, autonomous coding, and long-running agents [8]	General-purpose AI tasks, including coding and text/ image generation [3]
Context Window Size	Up to 1 million tokens [8]	Up to 200K or 1M tokens (for beta versions of Opus/ Sonnet) [3]
Agentic Capabilities	Explicitly built for advanced reasoning, planning, and execution by AI agents [1] [8]	General reasoning capabilities, with adaptive thinking in Sonnet/Opus [3]
Access Method	Gated research preview via Amazon Bedrock (`anthropic.claude-mythos-preview`) [8]	Standard API access via Anthropic Console or Bedrock [3]
Pricing	Information not available in provided sources	Opus: $5/in, 25/$out per million tokens; Sonnet: $3/in, 15/$out; Haiku: $1/in, 5/$out [3]

The Offensive Paradigm Shift: Mythos and the New Threat Landscape

The introduction of a frontier model like Claude's Mythos is poised to fundamentally alter the offensive cyber threat landscape, accelerating two convergent trends: the democratization of advanced attack capabilities and the industrialization of cyberattacks [2]. These shifts collectively lower the barrier to entry for conducting sophisticated cyber operations and increase the frequency and scale of attacks, placing immense pressure on existing defense mechanisms. The core concern is that the gap between machine-speed AI actions and human-speed security controls is becoming the primary source of risk for enterprises [7]. Mythos, with its purported capabilities in vulnerability discovery and exploit development, is at the center of this emerging reality [2]. Its ability to analyze software, identify vulnerabilities, and develop exploits presents a new and potent threat surface, making even legacy systems transparent to automated analysis [7].

The democratization of attack capabilities refers to the phenomenon where sophisticated techniques previously accessible only to elite, well-resourced adversaries, such as nation-

states, become available to a much broader range of actors, including low-skill individuals or less-resourced groups [2]. An AI system proficient in generating software could, with appropriate training or prompting, be turned toward identifying its own vulnerabilities [2]. When combined with Mythos's alleged multi-step attack reasoning, this creates a scenario where an individual with minimal cybersecurity expertise could orchestrate complex, multi-stage attacks [2]. This expansion of the potential adversary pool exponentially increases the overall threat volume and diversity. Furthermore, the convergence of this democratization with the industrialization of cyberattacks—a trend enabled by the automation and scalability of AI models—means that attackers can execute more sophisticated attacks simultaneously [2]. Instead of one expert crafting a single, bespoke exploit, an AI agent like Mythos could generate and deploy millions of variations of an attack vector across a target landscape in a matter of minutes, collapsing the traditional time-to-exploit window to near zero-day levels [2]. The example of Mythos reportedly discovering a 27-year-old flaw in OpenBSD serves as a stark illustration of how deeply and thoroughly these models can penetrate legacy codebases, rendering years of assumed security obsolete overnight [7].

This shift also alters the fundamental nature of the attack surface. As AI agents become more autonomous, the endpoint is no longer just a passive access point but an active execution environment for decision-making processes [7]. This evolution moves the security challenge beyond the perimeter and into the runtime governance of AI itself. The potential blast radius of a failure or malicious action by an uncontrolled AI agent is significantly larger than that of a traditional malware payload, as it can dynamically adapt its behavior and interact with multiple systems based on real-time feedback [7]. Securing AI, therefore, necessitates a new approach that extends beyond endpoints to include the protection of identity, context, and control planes that govern an agent's behavior [7]. The operational model of these agents represents a structural shift in enterprise security, demanding new methods to govern what they are allowed to do, under what conditions, and how their actions are continuously validated [7].

In response to these threats, the defensive landscape is being forced to evolve rapidly towards autonomy. The primary bottleneck in modern security is identified as the slow speed of human intervention, which cannot keep pace with the velocity of AI-driven attacks [7]. Consequently, the future of defense lies in developing autonomous systems capable of matching the speed and scale of the threats they face. This involves moving from reactive monitoring and signature-based detection to proactive vulnerability discovery and predictive threat hunting powered by AI [1]. For instance, Anthropic's Threat Intelligence team successfully discovered and disrupted a nation-state operation that utilized Claude’s agentic capabilities, demonstrating the potential for using similar AI

models to fight fire with fire [6]. By integrating frontier models like Mythos with real-world threat intelligence and enterprise-scale visibility, defenders can accelerate their ability to discover vulnerabilities, detect threats, and respond to incidents [1]. This proactive stance is critical, as the time available to patch a newly discovered vulnerability shrinks dramatically in an environment where exploit development is automated and instantaneous [2]. The strategic imperative for enterprises is to build defenses that are not just faster but are fundamentally designed to operate at machine speed, closing the dangerous gap between attack and response.

The Penetration Testing Debate: Obsolescence or Evolution?

The advent of a highly capable model like Claude's Mythos has ignited a fierce debate within the cybersecurity industry regarding the future viability of traditional penetration testing. The provocative question, "Is penetration testing dead?", encapsulates the central tension between the perceived efficiency of automated AI-driven scans and the irreplaceable value of human expertise. While some argue that the speed and analytical depth of AI will render manual testing obsolete, a more nuanced perspective suggests that penetration testing will not die but will instead undergo a significant evolution, shifting from a hands-on, labor-intensive practice to a higher-level strategic function focused on oversight, validation, and complex risk assessment. The reality likely lies somewhere between these two extremes, pointing toward a symbiotic relationship between human testers and AI-powered tools rather than outright replacement.

The argument for the obsolescence of traditional penetration testing is compelling and rests on the principles of automation, scale, and cost-efficiency. One of the primary drivers of this viewpoint is the economic model of vulnerability discovery. If large language models can find zero-day vulnerabilities at a fraction of the cost of human researchers, the perceived economic value of these findings could be drastically reduced [5]. With the cost of computation approaching zero, the incentive structure for bug bounty programs, which rely on rewarding human discovery, could collapse [5]. This suggests that any organization with access to a model like Mythos could perform continuous, automated scanning of its own or external targets, potentially finding and exploiting vulnerabilities faster and cheaper than a team of human pen testers working on a contract basis. Furthermore, the concept of "property-based fuzzing" being commodified by LLMs enables less-skilled individuals, sometimes derisively called "vibe-

hackers," to chain together millions of simple, previously unaddressed vulnerabilities to form complex, effective attacks with minimal investment [5]. This industrialized approach to exploitation bypasses the sequential, manual methodology that characterizes traditional penetration tests, which are often limited by budget, scope, and timeline constraints. In this view, the role of the human pen tester becomes redundant when an AI can perform the same task with greater speed, breadth, and persistence.

However, this perspective overlooks the enduring necessity of human judgment, creativity, and strategic context. Several sources indicate that while AI can automate the identification of technical flaws, it cannot fully replicate the multifaceted nature of a comprehensive security assessment. Penetration testing often extends beyond code to encompass physical security assessments, social engineering exercises, and evaluations of policy adherence and procedural compliance—all areas where human intuition and interaction are indispensable. An AI agent cannot convincingly impersonate an employee in a social engineering test or navigate the physical premises of a building. Moreover, the act of testing is not merely about finding a list of vulnerabilities; it is about understanding the business context of those vulnerabilities. A human expert is needed to assess the true business impact of a finding, prioritize risks based on asset value and regulatory requirements, and provide actionable, remediation advice tailored to the organization's specific environment and constraints. The unpredictable nature of generative AI makes human oversight essential for discerning true malicious intent versus benign anomalies or false positives that an algorithm might generate [7]. Therefore, the role of the penetration tester is likely to evolve from a hands-on scanner to a strategist who oversees AI-driven tools, validates their findings, interprets their outputs, and provides the final layer of human-centric analysis and recommendation.

Ultimately, the future of penetration testing will likely involve a hybrid model where AI augments, rather than replaces, human expertise. Executives in the security industry should anticipate a shift in the skill set required for their teams. There will be a growing demand for professionals who can manage and interpret the results of AI-driven security operations, focusing on higher-level risks, threat modeling, and strategic defense planning. The human element will remain critical for validating AI-generated findings, providing the contextual understanding that machines lack, and navigating the complex interplay between technology, policy, and human behavior. The demise of traditional, purely manual penetration testing is plausible, but it will be supplanted by a new discipline of AI-assisted security assessment, where human experts act as the crucial validators and strategists in an increasingly automated world. The key challenge for organizations will be to re-skill their security workforce to thrive in this new paradigm, ensuring that human oversight remains a cornerstone of their security posture.

The Defensive Imperative: Securing the Autonomous Era

In direct response to the escalating threat posed by autonomous AI agents like Mythos, the cybersecurity protection landscape is undergoing a necessary and profound evolution. The central challenge confronting defenders is the widening chasm between the speed of machine-driven attacks and the limitations of human-paced security controls [7]. To remain viable, enterprise security strategies must pivot from a reactive, signature-based model to a proactive, autonomous one capable of operating at machine speed. This transition demands a fundamental redesign of security architecture, a strategic embrace of AI-native defense technologies, and a re-evaluation of core security postures to address the expanded attack surface created by intelligent agents. The guiding principle for 2026 and beyond is clear: defenders must learn to fight fire with fire [6].

A primary pillar of this new defensive paradigm is the adoption of proactive vulnerability discovery. Traditional security relies heavily on patching known Common Vulnerabilities and Exposures (CVEs), leaving organizations vulnerable during the inevitable time lag between a vulnerability's disclosure and its remediation. Advanced AI models like Mythos can reverse this process, enabling defenders to discover and patch vulnerabilities before they are ever weaponized by an attacker [1]. By leveraging models with massive context windows, security teams can scan entire applications, networks, or even supply chains for weaknesses that would be impossible for humans to find manually [8]. The partnership between Anthropic and CrowdStrike exemplifies this forward-looking approach. By integrating the Mythos Preview with CrowdStrike's Falcon platform, which incorporates real-time threat intelligence and enterprise-scale visibility, defenders can create a powerful feedback loop: AI identifies potential issues, threat intelligence informs the context of those issues, and the platform enables rapid response [1]. This allows defenders to move from a state of constant catch-up to one of proactive discovery, detection, and response, compressing the attack lifecycle in their favor.

Beyond proactive discovery, the defensive imperative strongly points toward the development and deployment of autonomous response systems. The bottleneck in modern incident response is consistently identified as the delay caused by human involvement [7]. An AI-driven attack can compromise a system in seconds; a human analyst reviewing alerts and authorizing a response takes minutes or hours. This mismatch is untenable. The future of defense lies in systems that can automatically detect, contain, and remediate threats at machine speed. CrowdStrike's offerings, such as its AI Detection and Response (AIDR) technology, are designed to discover and govern all

AI agents at runtime, preventing malicious behavior before it can cause widespread damage [1]. Similarly, Falcon Data Security aims to prevent sensitive data from leaking out through compromised AI workflows, addressing a critical new vector of risk [1]. These technologies represent a shift from static rules and patterns to dynamic, intent-based analysis. Since generative AI is inherently unpredictable, relying on pattern matching is insufficient; defenses must evolve to understand the intent behind a request or action, a task better suited to AI-driven analysis [7].

To effectively implement these advanced defensive strategies, enterprises must conduct a rigorous reassessment of their security posture. Check Point's recommendations provide a practical framework for this adaptation, urging security leaders to focus on five key areas. First, organizations must evaluate their first-line defenses—such as firewalls, Web Application Firewalls (WAFs), and Endpoint Detection and Response (EDR) systems—for their ability to protect against zero-day attacks, a capability now demonstrably possible with AI. Second, vendors should be assessed based on their CVE history, not just as a measure of past performance but as a strategic liability that exposes the enterprise to future risk [2]. Third, security teams must actively hunt for blind spots in legacy servers and unpatched systems, which are prime targets for AI-driven scanners. Fourth, patching cycles must be accelerated to match the increased velocity of threat discovery. Finally, network segmentation must be reinforced to ensure that even if an attacker breaches one part of the network, they cannot easily move laterally to access critical assets [2]. The overarching conclusion is that securing the autonomous era requires a holistic architectural overhaul, prioritizing speed, automation, and the ability to govern intelligent agents at every stage of their lifecycle.

Financial Consequences and Market Implications

The emergence of advanced AI models like Claude's Mythos carries profound financial consequences for the cybersecurity industry, extending from stock market volatility and vendor business model disruption to enterprise risk management and the cyber insurance sector. The mere leak of information regarding the model's capabilities was sufficient to trigger a sharp sell-off in cybersecurity stocks, signaling that financial markets perceive this technology as a significant disruptive force that could reshape the competitive landscape and erode the value of incumbent players whose services may become commoditized [4]. This immediate market reaction highlights investor anxiety about the potential for AI to automate tasks currently performed by humans, thereby reducing the demand for traditional security services and compressing profit margins.

One of the most immediate financial impacts will be on the business models of security vendors. The widespread adoption of LLM-based security scanning could force software vendors, such as Cisco, ServiceNow, AWS, and Microsoft, to integrate these capabilities into their core products as a standard feature rather than a separate service [5]. This would effectively absorb the cost of the technology, turning it from a revenue-generating module into a cost center, which could lead to margin erosion for pure-play security companies that do not have a similar scale advantage. Furthermore, the economic foundation of bug bounty programs—a significant revenue stream for independent researchers and a valuable source of intelligence for companies—could be undermined. If the cost to find vulnerabilities approaches zero due to cheap, automated scanning, the incentive for skilled human researchers to participate in these programs diminishes, potentially disrupting a key component of the vulnerability disclosure ecosystem [5]. This shift forces security firms to innovate and find new value propositions beyond basic vulnerability scanning, perhaps focusing on the strategic analysis and remediation guidance that AI alone cannot provide.

The increased velocity and sophistication of AI-powered attacks will inevitably have a cascading effect on enterprise risk management and the cyber insurance industry. Insurers are likely to see a rise in both the frequency and severity of claims as automated attacks become more prevalent and damaging [1]. In response, insurers may tighten their underwriting criteria, making coverage more difficult and expensive to obtain. A potential outcome is the inclusion of clauses in cyber insurance contracts that mandate the use of LLM-based scans as a condition of coverage, effectively forcing enterprises to adopt the very technology that poses a threat [5]. This creates a paradoxical situation where the solution to managing AI-driven risk becomes a contractual obligation tied to financial protection. Similarly, government and large corporate procurement processes may begin to mandate strict remediation timelines enforced by such automated scanning technologies, compelling vendors to invest heavily in AI-native security tools to remain competitive [5]. Companies that fail to adapt their security posture to this new reality may find themselves facing higher insurance premiums, stricter contractual obligations, and increased exposure to catastrophic breaches that could have devastating financial repercussions.

Finally, the proliferation of powerful AI models raises important questions about sustainability and long-term costs. There are significant concerns about the environmental and energy grid impacts associated with running these computationally intensive systems at scale [5]. The substantial energy consumption required for training and operating frontier models like Mythos introduces a societal cost-benefit consideration that could influence public perception, corporate governance policies, and even future

regulation. While the immediate financial impact is dominated by market reactions and business model pressures, the long-term economic viability of these technologies will depend on their ability to balance immense power with sustainable operational costs. For executives, this means that the financial implications of adopting AI in security extend beyond the initial capital expenditure to include ongoing operational costs, insurance liabilities, and the strategic positioning required to maintain a competitive edge in a rapidly evolving and financially turbulent market.

Strategic Recommendations for Executive Leadership

The introduction of Anthropic's Mythos model signals a pivotal moment for the cybersecurity industry, demanding a strategic recalibration from enterprise leadership. The confluence of AI-driven threats and the need for autonomous defense systems renders outdated many traditional security paradigms. To navigate this new landscape successfully, executives must move beyond incremental adjustments and undertake a comprehensive review of their security posture, business models, and strategic partnerships. The following recommendations provide a roadmap for adapting to the realities of an AI-powered era, focusing on proactive defense, human-AI collaboration, and long-term resilience.

First, executives must initiate an immediate and thorough audit of their current security stack to identify vulnerabilities in processes and technologies that are slow, manual, or predicated on outdated assumptions. This audit should specifically target areas where the speed of response is a critical factor, such as incident detection and containment. The strategic imperative is to prioritize investment in AI-native security solutions that offer automation, high velocity, and autonomous response capabilities [7]. This includes evaluating technologies like AI Detection and Response (AIDR) and other platforms designed to govern and react to AI agent behavior at runtime [1]. The goal is to close the dangerous gap between machine-speed AI actions and human-speed security controls by embedding intelligence and automation directly into the core of the security infrastructure [7]. Relying on legacy tools and manual processes will become increasingly untenable as the threat landscape accelerates.

Second, organizations must foster a deliberate and structured "human-in-the-loop" strategy rather than attempting to replace human expertise entirely. The evidence suggests that the future of security lies in a symbiotic relationship between AI and humans, not in the displacement of one by the other [7]. Human oversight will remain

critical for interpreting the outputs of AI-driven security operations, validating findings, and applying contextual business judgment that algorithms lack. Executives should invest in re-skilling their security teams to operate as strategists, validators, and analysts for AI-driven systems. This involves developing clear protocols for how humans will oversee automated decisions, intervene in critical situations, and leverage AI-generated insights to inform higher-level risk management and strategic planning. The role of the human expert will shift from performing tactical tasks to providing the essential final layer of intelligent oversight.

Third, executive leadership should proactively engage in scenario planning and red-teaming exercises that specifically simulate attacks orchestrated by autonomous AI agents like Mythos. Creating realistic attack scenarios helps to illuminate architectural weaknesses and reveals gaps in incident response plans that may not be apparent under conventional threat models. These simulations should explore how an AI adversary could breach defenses, bypass controls, escalate privileges, and exfiltrate data. The insights gained from these exercises are invaluable for informing architectural changes, strengthening network segmentation, and refining response playbooks to be effective against threats that operate at machine speed [2]. This proactive approach to stress-testing the security posture is far more effective than waiting for a breach to expose systemic vulnerabilities.

Fourth, it is imperative to critically review existing vendor and partnership agreements through the lens of the new AI-driven threat landscape. Executives should assess their relationships with security vendors, seeking partners who are demonstrably integrating AI capabilities into their products and have a transparent roadmap for AI security [1] [7]. Contractual discussions should incorporate clauses related to AI-assisted vulnerability discovery, remediation timelines, and the sharing of threat intelligence. Forming strategic alliances with firms at the forefront of AI security, such as Anthropic's partnership with CrowdStrike, can provide early access to advanced defensive capabilities and valuable threat intelligence [1]. A vendor's commitment to leveraging AI for defense should become a key criterion in procurement decisions.

Finally, leaders must maintain a vigilant watch on the rapidly evolving competitive and regulatory landscape. The cybersecurity market is poised for a shake-up, where competitors who successfully leverage AI for defense will gain a significant strategic advantage [4]. This could impact customer acquisition, retention, and brand reputation. Simultaneously, executives should monitor for emerging regulations that may mandate the use of AI-based scanning for certain industries or government contractors, as well as shifts in cyber insurance requirements that could make such technologies de facto

standards [5]. By staying ahead of these external pressures, organizations can avoid being caught off guard and can position themselves as leaders in the new era of autonomous security.

Reference

1. Anthropic Claude Mythos Preview - CrowdStrike https://www.crowdstrike.com/en-us/blog/crowdstrike-founding-member-anthropic-mythos-frontier-model-to-secure-ai/
2. Claude Mythos wake-up call: What AI vulnerability discovery means ... https://www.intelligentciso.com/2026/04/09/claude-mythos-wake-up-call-what-ai-vulnerability-discovery-means-for-cyberdefence/
3. Models overview - Claude API Docs https://docs.anthropic.com/en/docs/about-claude/models/all-models
4. AI leak triggers sharp cybersecurity sell-off - MSN https://www.msn.com/en-us/news/insight/ai-leak-triggers-sharp-cybersecurity-sell-off/gm-GM32B19C6B?gemSnapshotKey=GM32B19C6B-snapshot-6
5. Some quick thoughts on Claude Mythos and the economics of ... https://www.linkedin.com/posts/malwaretech_some-quick-thoughts-on-claude-mythos-and-activity-7448887517703188480-lHFM
6. Author - CrowdStrike.com https://www.crowdstrike.com/en-us/blog/author.crowdstrike/
7. George Kurtz's Post - Anthropic Claude Mythos Preview - LinkedIn https://www.linkedin.com/posts/georgekurtz_anthropic-claude-mythos-preview-the-more-activity-7447349706000826368-OG3E
8. Claude Mythos Preview - Amazon Bedrock - AWS Documentation https://docs.aws.amazon.com/bedrock/latest/userguide/model-card-anthropic-claude-mythos-preview.html